

Interview - Codility

app.codility.com/c/interview/CL-FH3QYE-X8Z/#/

BookmarksoldMapscode allML allsystemfinance allpersonal allmathtooltmp

ASKS

1 > Task 1

SYSTEM DESIGN

Whiteboard

Take a tour

Accessibility mode

Settings

Task 1

Simulated Distributed ML Training with Checkpointing

Objective:

Implement a simplified version of a distributed machine learning (ML) training job to demonstrate your understanding of parallelism, data partitioning, and fault tolerance.

Problem Statement:

You are tasked with simulating a distributed ML training job using multiple parallel trainer processes. The training is simplified to summing numeric data.

Requirements:

- Input Data:**
 - Use a CSV file or an in-memory array that contains one numeric value per row (e.g., `[1, 2, 3, ..., 1000]`).
- Parallel Trainers:**
 - Launch `N` parallel "trainer" processes.
 - Each trainer should load and process a distinct shard of the data (non-overlapping rows).
- Simulated Training:**
 - Each trainer computes the sum of its assigned shard.
 - Then, combine the partial results to compute the global sum.
- (Bonus) Checkpointing Support:**
 - Implement a mechanism to save the intermediate state (e.g., each trainer's partial sum) to disk.
 - Provide functionality to resume training from the last checkpoint in the event of a failure.
- Testing:**
 - Test the implementation under:


Feedback

EXPLORE

WORKSPACE

OUTLINE


TIMELINE





TASKS

< 1 > Task 1


SYSTEM DESIGN

 Whiteboard

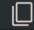
 Take a tour

 Accessibility mode

Task 1

- Use a CSV file or an in-memory array that contains one numeric value per row (e.g., `[1, 2, 3, ..., 1000]` ).

2. Parallel Trainers:

- Launch `N`  parallel "trainer" processes.
- Each trainer should load and process a distinct shard of the data (non-overlapping rows).

3. Simulated Training:

- Each trainer computes the sum of its assigned shard.
- Then, combine the partial results to compute the global sum.

4. (Bonus) Checkpointing Support:

- Implement a mechanism to save the intermediate state (e.g., each trainer's partial sum) to disk.
- Provide functionality to resume training from the last checkpoint in the event of a failure.

5. Testing:

- Test the implementation under:
 - A normal case (no failure): verify the total sum matches the expected result.
 - (Bonus) A failure-recovery case: simulate a process failure, resume from checkpoint, and verify that the final result is still correct.

Task 1

Simulated Distributed ML Training with Checkpointing

Objective:
Implement a simplified version of a distributed machine learning (ML) training job to demonstrate your understanding of parallelism, data partitioning, and fault tolerance.

Problem Statement:
You are tasked with simulating a distributed ML training job using multiple parallel trainer processes. The training is simplified to summing numeric data.

Requirements:

1. **Input Data:**
 - Use a CSV file or an in-memory array that contains one numeric value per row (e.g., `[1, 2, 3, ..., 1000]`).
2. **Parallel Trainers:**
 - Launch `N` parallel "trainer" processes.
 - Each trainer should load and process a distinct shard of the data (non-overlapping rows).
3. **Simulated Training:**
 - Each trainer computes the sum of its assigned shard.
 - Then, combine the partial results to compute the global sum.
4. **(Bonus) Checkpointing Support:**
 - Implement a mechanism to save the intermediate state (e.g., each trainer's partial sum) to disk.
 - Provide functionality to resume training from the last checkpoint in the event of a failure.
5. **Testing:**
 - Test the implementation under:
 - A normal case (no failure): verify the total sum matches the expected result.
 - (Bonus) A failure-recovery case: simulate a process failure, resume from checkpoint, and verify that the final result is still correct.

Feedback

EXPLORER

WORKSPACE

main.py 1

main.py 1

```
1 """ Using any AI assist is prohibited. """
2
3 import multiprocessing as mp
4
5
6 def worker(proc_id, shard, results):
7     results[proc_id] = 0
8
9
10 def main():
11     manager = mp.Manager()
12     results = manager.list([None])
13     p = mp.Process(target=worker, args=(0, [], results))
14     p.start()
15     p.join()
16     results = list(results)
17     print(results)
18
19
20 if __name__ == "__main__":
21     main()
22
23 # Zhefu? Can you
```

Task 1

Simulated Distributed ML Training with Checkpointing

Objective:
Implement a simplified version of a distributed machine learning (ML) training job to demonstrate your understanding of parallelism, data partitioning, and fault tolerance.

Problem Statement:
You are tasked with simulating a distributed ML training job using multiple parallel trainer processes. The training is simplified to summing numeric data.

Requirements:

- Input Data:**
 - Use a CSV file or an in-memory array that contains one numeric value per row (e.g., `[1, 2, 3, ..., 1000]`).
- Parallel Trainers:**
 - Launch `N` parallel "trainer" processes.
 - Each trainer should load and process a distinct shard of the data (non-overlapping rows).
- Simulated Training:**
 - Each trainer computes the sum of its assigned shard.
 - Then, combine the partial results to compute the global sum.
- (Bonus) Checkpointing Support:**
 - Implement a mechanism to save the intermediate state (e.g., each trainer's partial sum) to disk.
 - Provide functionality to resume training from the last checkpoint in the event of a failure.
- Testing:**
 - Test the implementation under:
 - A normal case (no failure): verify the total sum matches the expected result.
 - (Bonus) A failure-recovery case: simulate a process failure, resume from checkpoint, and verify that the final result is still correct.

Feedback

main.py 1

```
1 """ Using any AI assist is prohibited. """
2
3 import multiprocessing as mp
4
5
6 def worker(proc_id, shard, results):
7     results[proc_id] = 0
8
9
10 def main():
11     manager = mp.Manager()
12     results = manager.list([None])
13     p = mp.Process(target=worker, args=(0, [], results))
14     p.start()
15     p.join()
16     results = list(results)
17     print(results)
18
19
20 if __name__ == "__main__":
21     main()
22
23 # Zhifu? Can you
```

OUTLINE

```
//-----  
myV0
```

Problem Statement:

You are tasked with simulating a distributed ML training job using multiple parallel trainer processes. The training is simplified to summing numeric data.

Requirements:

1. Input Data:

- Use a CSV file or an in-memory array that contains one numeric value per row (e.g., `[1, 2, 3, ..., 1000]`).

2. Parallel Trainers:

- Launch `N` parallel "trainer" processes.
- Each trainer should load and process a distinct shard of the data (non-overlapping rows).

3. Simulated Training:

- Each trainer computes the sum of its assigned shard.
- Then, combine the partial results to compute the global sum.

Millie stones:

- Sum all the numbers from 1 to 1000 and print out using `N` processes (say 10 processes)
- Sum of the number at every training step
- Load the data data in random order
- Make the order (still random) deterministic between 2 training runs

Task 1

Problem Statement:

You are tasked with simulating a distributed ML training job using multiple parallel trainer processes. The training is simplified to summing numeric data.

Requirements:

- Input Data:**
 - Use a CSV file or an in-memory array that contains one numeric value per row (e.g. `[1, 2, 3, ..., 1000]`).
- Parallel Trainers:**
 - Launch `N` parallel "trainer" processes.
 - Each trainer should load and process a distinct shard of the data (non-overlapping rows).
- Simulated Training:**
 - Each trainer computes the sum of its assigned shard.
 - Then, combine the partial results to compute the global sum.

Milieu stones:

- Sum all the numbers from 1 to 1000 and print out using `N` processes (say 10 processes)
- Sum of the number at every training step
- Load the data data in random order
- Make the order (still random) deterministic between 2 training runs

Feedback

EXPLORER

WORKSPACE

main.py 1

```
1 import multiprocessing as mp
2
3 # you can write to stdout for debugging purposes, e.g.
4 print("This is a debug message")
5
6
7 ## 10 processes
8 ##
9
10 DATA_SIZE = 1000
11 NUM_PROC = 4
12
13
14 # generate training data for each worker
15 def createDataPart(data, num_proc):
16     average_len = len(data) // num_proc
17     data_parts = []
18     for idx in range(num_proc):
19         start = idx * average_len
20         end = (idx+1) * average_len
21         if idx == num_proc - 1:
22             end = len(data)
23
24         tup = data[start:end]
25         data_parts.append(tup)
26
27     return data_parts
28
29
30 # each worker logic
31 def worker(proc_id, data_part, results):
32     local_sum = 0
33     start=0
34
35     len0 = len(data_part)
36     for idx in range(len0):
37         local_sum += data_part[idx]
38
39     results[proc_id] = local_sum
40     print("proc_id = ", proc_id, ", local_sum = ", local_sum)
41     return
42
43 # simulate the parallel training
44 def simulation(data, num_proc):
45     data_parts = createDataPart(data, num_proc)
46
47     manager = mp.Manager()
```

Task 1

Problem Statement:

You are tasked with simulating a distributed ML training job using multiple parallel trainer processes. The training is simplified to summing numeric data.

Requirements:

- Input Data:**
 - Use a CSV file or an in-memory array that contains one numeric value per row (e.g., `[1, 2, 3, ..., 1000]`).
- Parallel Trainers:**
 - Launch `N` parallel "trainer" processes.
 - Each trainer should load and process a distinct shard of the data (non-overlapping rows).
- Simulated Training:**
 - Each trainer computes the sum of its assigned shard.
 - Then, combine the partial results to compute the global sum.

Millie stones:

- Sum all the numbers from 1 to 1000 and print out using `N` processes (say 10 processes)
- Sum of the number at every training step
- Load the data data in random order
- Make the order (still random) deterministic between 2 training runs

Feedback

main.py 1

```
31 def worker(proc_id, data_part, results):
32     .....
34
35     len0 = len(data_part)
36     for idx in range(len0):
37         local_sum += data_part[idx]
38
39     results[proc_id] = local_sum
40     print("proc_id = ", proc_id, ", local_sum = ", local_sum)
41     return
42
43 # simulate the parallel training
44 def simulation(data, num_proc):
45     data_parts = createDataPart(data, num_proc)
46
47     manager = mp.Manager()
48     results = manager.list([None] * num_proc)
49     procs = []
50
51     for idx in range(num_proc):
52         p_arg = (idx, data_parts[idx], results)
53
54         p = mp.Process(target=worker, args=p_arg)
55         procs.append(p)
56         p.start()
57
58     for p in procs:
59         p.join()
60
61     rst = 0
62     for tmp in results:
63         if tmp != None:
64             rst += tmp
65
66     results = list(results)
67     print(results)
68
69     print("rst = ", rst)
70     return rst
71
72
73 def main():
74     data = list(range(1, 1+DATA_SIZE))
75     expected_rst = sum(data)
76
77     final_sum = simulation(data, NUM_PROC)
78     print("expected_rst = ", expected_rst, ", final_sum = ", final_sum)
79
```

#-----
simulating distributed ml training with checkpointing

objective:

implement a simplified version of a distributed machine learning (ml) training job to demonstrate your understanding of parallelism, data partitioning, and fault tolerance

problem statement:

you are tasked with simulating a distributed ml training job using parallel trainer processes. the training is simplified to summing numeric data

requirements:

1. input data

using a csv file or an in-memory array that contains one numeric value per row (e.g., [1,2,3,..., 1000])

2. parallel trainers:

launch N parallel "trainer" processes

each trainer should load and process a distinct shard of the data (non-overlapping rows)

3. simulated training

each trainer computes the sum of its assigned shard

then combine the partial results to compute the global sum

4. sum of the number at every training step

5. load the data in random order

6. make the order (still random) deterministic between 2 training runs

7. please use this code as an example

