

pytorch attention implementation code, with both forward propagation, and backward propagation code, please explain with examples

💡 Thought for 1m 40s

## Custom Scaled Dot-Product Attention in PyTorch with Forward and Backward Propagation

In PyTorch, the attention mechanism (particularly scaled dot-product attention) is typically handled by built-in functions like

`torch.nn.functional.scaled_dot_product_attention` or `nn.MultiheadAttention`, which rely on autograd for backward propagation. However, to explicitly show both forward and backward passes, we'll implement a custom version using `torch.autograd.Function`. This allows us to define the forward computation and manually compute gradients in the backward pass for educational purposes.

This implementation is for **single-head** scaled dot-product attention. For multi-head, you can parallelize it across heads. We'll ignore masks and dropout for simplicity (they can be added similarly to the previous response).

### Derivation Recap

- **Forward Pass:**
  - Scores  $S = Q \cdot K^T \times \text{scale}$ , where  $\text{scale} = 1/\sqrt{d_k}$ .
  - Attention weights  $A = \text{softmax}(S)$ .
  - Output  $O = A \cdot V$ .
- **Backward Pass (Gradients):**

you can parallelize it across heads. We'll ignore masks and dropout for simplicity (they can be added similarly to the previous response).

### Derivation Recap

- **Forward Pass:**
  - Scores  $S = Q \cdot K^T \times \text{scale}$ , where  $\text{scale} = 1/\sqrt{d_k}$ .
  - Attention weights  $A = \text{softmax}(S)$ .
  - Output  $O = A \cdot V$ .
- **Backward Pass (Gradients):**
  - Given incoming gradient  $\frac{\partial L}{\partial O}$  (grad\_output).
  - $\frac{\partial L}{\partial V} = A^T \cdot \frac{\partial L}{\partial O}$ .
  - $\frac{\partial L}{\partial A} = \frac{\partial L}{\partial O} \cdot V^T$ .
  - Softmax backward:  $\frac{\partial L}{\partial S} = A \odot \frac{\partial L}{\partial A} - A \odot (A^T \cdot \frac{\partial L}{\partial A})$  (more precisely, row-wise:  $A * \frac{\partial L}{\partial A} - A * (A * \frac{\partial L}{\partial A}).\text{sum}(\text{dim} = -1, \text{keepdim} = \text{True})$ ).
  - $\frac{\partial L}{\partial Q} = \frac{\partial L}{\partial S} \cdot K \times \text{scale}$ .
  - $\frac{\partial L}{\partial K} = Q^T \cdot \frac{\partial L}{\partial S} \times \text{scale}$ .

The code below implements this exactly.

### Code Implementation

python

✕ Collapse

≡ Wrap

▶ Run

📄 Copy

```
import torch
import math
```

